

SUPPLEMENTARY MATERIAL

OPEN ACCESS

Supplementary Material: Agentic Autoresearch for CT Reconstruction

 Andreas Maier¹ | Lukas Kachelrieß^{1,2} | Siming Bayer¹ | Yixing Huang³ | Yan Xia² | Amber Simpson⁴ | Moritz Zaiss²
¹Pattern Recognition Lab, Friedrich-Alexander-Universität Erlangen-Nürnberg (FAU), Erlangen, Germany | ²Universitätsklinikum Erlangen (UKER), Erlangen, Germany | ³Peking University, Beijing, China | ⁴University of Alberta, Edmonton, Canada

Correspondence: Andreas Maier (andreas.maier@fau.de)

Received: 9 July 2026 | **Revised:** 9 July 2026 | **Accepted:** 9 July 2026

Keywords: CT reconstruction | sparse-view | low-dose | known operators | benchmarking | noise robustness

ABSTRACT

This document contains the supplementary material for *Agentic Autoresearch for CT Reconstruction*. It provides the full solver taxonomy and per-method configuration (S1), the Mayo-LDCT data, split, and geometry pipeline (S2), the breast split and train/test leakage audit (S3), the Mayo test board and significance (S4), the parameter-efficient recombination study (S5), the noise model and no-retrain protocol (S6; the full boards and the framework-axis reading are in the main text), an extended account of agent behavior (S7), and additional results with pointers to the public repository and live leaderboard (S8).

S1 | Solver taxonomy and per-method configuration (Table 1, full)

We place the 26 methods on two axes. Axis A is physics engagement (how the operator **A** is used at inference). Axis B is the prior source $\mathcal{R}$.

Full per-solver CONFIG defaults, parameter counts, and design notes are in each solver's design doc under `pentathlon/demo_dl_reference/` and the observation records under `docs/runs/`.

S2 | Mayo-LDCT — data, split, and the geometry pipeline

Real helical low-dose CT from the 2016 AAPM Low-Dose CT Grand Challenge.²³ We rebinned helical to fan-beam by single-slice rebinning.²⁴ The differentiable fan-beam projector used at every step is PYRO-NN.²⁵ This pipeline was the dominant human cost of the whole study (~24 active working days; three travel gaps excluded). We use a fixed patient split: train L145/186/209/219, validation L277, test L014/056/058/075/123. Test scores are the per-patient mean $\pm$ std over the five test patients ($n = 5$). Geometry details, rebinning validation, and the effort timeline figure are in the repository (`docs/runs/mayo_effort_timeline.png`).

The Mayo challenge is one of a family of AAPM reconstruction grand challenges spanning sparse-view, spectral, truth-based, and metal-artifact tasks.^{26,27,28,29,30}

S3 | Breast split and the train/test leakage audit

Split. We re-partitioned the public 4000-case Breast CT Challenge train pool into train (3600), validation (200), and test (200). The split is deterministic (seed 20260703) and disjoint. The official challenge val/test truths are withheld, so we cannot use them; our test split is drawn from the public pool.

Leakage audit (three independent checks).

1. *Mechanism.* Training loads the train split. The test redirect fires only for the evaluation load. Training never sees test.
2. *Pixel-hash disjointness.* Over all 3600 train, 200 val, and 200 test truth images: $\text{TEST} \cap \text{TRAIN} = 0$, $\text{TEST} \cap \text{VAL} = 0$, $\text{VAL} \cap \text{TRAIN} = 0$.
3. *Behavioral.* Test-hr tracks val-hr for every method, within ± 0.01 , and several methods score *lower* on test than val. Training on test would make test-hr much higher. It does not. The high absolute hr is the dataset (noiseless, well-posed), not a leak.

Why breast hr is high, and not comparable to Mayo. The Breast CT Challenge data is noiseless by design; its RMSE floor

TABLE S1 | Full solver taxonomy. Axis A = physics engagement; Axis B = prior source. The final column cites the source(s) for each family. Abbreviations: DC, data consistency; N2I, Noise2Inverse; rectified-flow prior, a fast diffusion-type generative prior; FoE, field of experts (a learned filter-bank image prior).

Family	Representative solver(s)	Data-consistency D	Prior $\mathcal{R}$	Axis A	Axis B	Ref.
Classical iterative	tv-iterative, manhart-PWLS-TV	filtered / SART grad	total variation	in-loop	hand-crafted	1,2
Image-domain denoiser	manduca-bilateral, DD-UNet sup.	none (post-FBP)	bilateral / U-Net	none	hand-cr. / sup.	3,4
Unrolled learned-iterative	ITNet v1/v2/v3, U-Swin	in-loop DC	learned CNN / transformer	in-loop	supervised	5,6,7,8,9
Variational network	Hammernik-2017, Hammernik-VN	in-loop DC	field-of-experts	in-loop	supervised	10,11
Learned primal-dual	learned-primal-dual (LPD)	filtered + learned dual	learned combination	in-loop	supervised	12
Dual-domain	DD-supervised, DD-bilateral, DD-N2I	image + sinogram	sup. / bilateral / self-sup.	single/in-loop	sup. / self-sup.	13,14
Generative prior	fastdiff flow/wavelet $\times 4$	DC + diffusion	rectified-flow prior	in-loop	generative	15,16,17,18
Per-scene implicit	NAF, R^2 -Gaussian	per-scene fit	implicit representation	per-scene	implicit	19,20,21
Foundation zero-shot	RAM	single DC	frozen foundation model	single	foundation	22
Band decomposition	Wu-2015, Wu-2015-trainable	band-weighted DC	frequency-band residual	in-loop	hand-cr. / sup.	2
Recombination (ours)	param-efficient	filtered + LPD combine	bilateral post-filter	in-loop	hybrid	—

TABLE S2 | Breast versus Mayo: the two problems are limited by different things, so their absolute hr scales are not comparable.

	Breast CT	Mayo
data	noiseless, ideal	real low-dose, quantum noise
bottleneck	data incompleteness	noise / dose
RMSE floor	0	> 0
best hr	~0.89	~0.38
dominant lever	data-consistency / known operator	denoising

is zero. Breast is incompleteness-limited (128 of ~1000 views). Mayo is noise-limited (real quantum noise). The two absolute hr scales are therefore not comparable.

S4 | Mayo test board and significance

Champion ITNet v1, test hr 0.376 ($n = 5$). The top tier is a statistical tie: ITNet v1 / ITNet v2 / U-Swin at the 5% level; the compact param-efficient solver joins at the 1% level ($p = 0.02$ – 0.027). At $n = 5$ the paired t -test has low power (a Wilcoxon test cannot even reach $p < 0.05$). The full board (all 26 solvers, hr/SSIM/RMSE) is in the main text; per-patient hr and the pairwise significance matrix are in docs/runs/mayo_significance_stats.md and the forest/matrix figures.

S5 | The parameter-efficient recombination study (search arc)

The agent recombined the method parts into a minimal-parameter solver on each dataset. Each climb was made of genuine six-box iterations, all under the 20-min budget, and the optimal recombination differs by problem. We present the Mayo arc first, then breast.

Mayo (low-dose, denoising). The seed is a learnable CNN regulariser in a tied prox-DC unroll. Switching the regulariser *family* from a CNN to a field-of-experts (FoE) filter bank set the agent-alone ceiling at hr ≈ 0.25 (Table S3). The break came from the projection domain: prepending a noise-adaptive Manduca bilateral filter to the sinogram, then extending it to a multi-scale bank, lifted val hr to 0.40. Halving the FoE bank yields a 969-parameter Pareto pick (test hr 0.324) — the reported Mayo compact solver.

TABLE S3 | Parameter-efficient recombination arc on Mayo (low-dose, denoising). hr is validation headroom; the last row is the reported low-parameter pick.

step	architecture	params	val hr
CNN prox+DC (seed)	learnable CNN regulariser	~2800	0.24
FoE filter bank	field-of-experts reg (agent ceiling)	1921	0.25 (ceiling)
+ Manduca projection filter	noise-adaptive sinogram bilateral	1926	0.394
+ multi-scale projection bank	parallel Manduca scales	1937	0.400 (best)
halved FoE (Pareto pick)	$n_f 24 \rightarrow 12$	969	0.360 (test 0.324)

TABLE S4 | Parameter-efficient recombination arc on breast (128-view sparse), all under the 20-min budget.

step	architecture	params	val hr
image-domain unroll (any config)	Mayo-style FoE denoiser	~1930	≤ 0.26 (ceiling)
filtered DC FBP($Ax - g$)	$K = 30$	13	0.4345
+ learned primal-dual combination	$W = 4$	183	0.5047
+ 3-scale bilateral output filter	$\sigma_r = 0.02$	195	0.6201 (single best)

Breast (128-view sparse). Key steps, all under the 20-min budget:

Findings, each earned by an iteration: the image-domain denoiser caps at ~ 0.26 ; deeper *adjoint* data-consistency lowers hr (it re-injects streaks); the *filtered* gradient breaks the ceiling at 13 parameters; an in-loop prox (proximal / denoising step) hurts (the filtered gradient already reconstructs); a learned cross-step combination helps; a full conjugate-gradient DC solve does not beat the single filtered gradient under a fixed wall; and Wu-2015 band machinery adds nothing here. The champion is seed-fragile in the variational-network manner: 4 of 5 seeds cluster at 0.571 ± 0.033 , one collapses; two stabilizers did not recover it. On the held-out test set the compact solver scores 0.6212 ± 0.0076 — the tightest per-case std of any solver.

The gains are back-loaded, and human-cued — on both datasets. They are not spread evenly across the ~40-iteration searches. Under the agent alone the best value plateaued early — at hr ≈ 0.25 on Mayo (by iteration 7) and ≈ 0.50 on breast (by iteration 15) — and then sat flat for a dozen or more iterations. Only in the second half, after a human began suggesting untried directions, did each break through: Mayo to hr 0.40 from iteration 25, breast to hr 0.62 at iteration 28 (Figure S1). This is the clearest evidence in the study of the agent needing human steering to escape a local optimum; without the hints it would have reported both compact solvers at their plateaus.

Mayo vs breast contrast. On Mayo the compact optimum is an evolved field-of-experts plus multi-scale bilateral denoiser (hr 0.324 at ~ 0.00097 M params), statistically tied with the top tier at the 1% level. On a denoising problem, a tiny denoiser suffices. On sparse-view breast, that same denoiser fails; data-consistency is the differentiator.¹² This problem-dependence is the practical face of the known-operator principle^{31,32}: embedding the exact forward operator cannot raise, and usually lowers, the maximum error bound, and it curbs the hallucination risk of unconstrained learned maps.^{33,34} The agent re-derived each optimum; it did not transfer.

Human role in this study. The recombination *strategy* was a human idea, on both datasets. On Mayo the human repeatedly asked whether more parameter-efficient approaches existed and explicitly suggested combining the pieces. On breast the finale was framed as a cross-method recombination and steered away from




FIGURE S1 | Compact-solver search trajectories. Top: Mayo (low-dose). Bottom: breast (128-view sparse). Points are per-iteration validation headroom (discarded iterations shown at hr 0); the step line is the best value so far. On both datasets the agent's own first 20 iterations plateau (hr ≈ 0.25 Mayo, ≈ 0.50 breast); the decisive break comes only in the hinted second half (shaded) — to hr 0.40 and 0.62 respectively.

copying Mayo. The agent executed and mechanistically refined; it did not originate the recombination strategy.

S6 | Noise model and no-retrain protocol

The full noiseless and noisy breast boards, and the framework-axis reading of the reversal, are now in the main text; this section records only the exact noise protocol behind the re-evaluation. Poisson photon noise per Eq. (6) of the main text, $I_0 = 10^5$, applied to the 200 test sinograms with a fixed seed (deterministic; every method sees the identical noisy input). RMS of the sinogram perturbation ≈ 0.012 (line-integral units), $\sim 1.8\%$ at the thickest ray. Supervised methods load their noiseless checkpoint and skip training; per-scene/classical methods re-fit on the noisy input. The FBP baseline is recomputed from the noisy sinogram, so hr measures improvement over the noisy FBP and the comparison is fair within the noisy board.




FIGURE S2 | Effect size on the breast board ($n = 200$): Cohen's d_z versus raw Δ hr against the champion, one point per solver. Large $|d_z|$ at $n = 200$ is why the breast board separates cleanly, unlike the $n = 5$ Mayo board.

S7 | Agent-behavior — extended account

This account follows the autoresearch paradigm of automated open-ended scientific discovery³⁵ realized through LLM code agents that resolve real software tasks^{36,37} and physics-informed agentic development in adjacent imaging fields.³⁸

Strengths. Fast paper→code; strong hyper-parameter optimization; strict fixed-budget discipline; scales evaluation across many parallel benchmarks; follows well-decomposed instructions.

Weaknesses. Does not invent new methods. Must be forced to recombine (converges early and pads if unsupervised — one breast sub-run replayed identical configs until an audit caught it). CT-image vision misses obvious artifacts, so numbers are the source of truth. Long tasks must be decomposed or even a 1M-token context is exhausted. Overfits to the task/metric (not only an agent trait).

Human meta-strategies (repeatable, same on both datasets). Persistence/breadth (keep asking for more directions); recombination (the idea to assemble pieces); redirection (do not mimic the other dataset); auditing (catch padding, enforce budget, fix provenance). The defensible claim: agent = tireless, mechanistically reasoning, self-modifying executor; human = strategist + auditor.

S8 | Additional results and the live leaderboard

One further figure is included here. Figure S2 shows, on the breast board ($n = 200$), the effect size (Cohen's d_z) against the raw Δ hr to the champion for each solver; with $n = 200$ even small hr gaps are highly significant, in contrast to the low-power $n = 5$ Mayo setting. The parameter/accuracy trade-off itself — for both datasets — is Figure 2 of the main text, and the per-step climb is quantified in the recombination-arc tables above (Tables S3 and S4).

Everything that did not fit the paper is public. The complete results live in the repository and the live dashboard:

<https://github.com/akmaier/Agent4CT>
<https://akmaier.github.io/Agent4CT/>

The dashboard hosts the full leaderboards for all datasets — Mayo, noiseless breast, and the noisy re-evaluation — with every solver’s per-iteration trajectory. Each iteration carries an immutable provenance record: the exact configuration, the reconstruction, and the frozen metric, so any number in this paper can be traced to the run that produced it. The repository additionally holds the Mayo significance forest plot and solver×metric matrix, the val-vs-test selection-honesty curves, the effort-timeline figure, the noiseless-versus-noisy reconstruction montages for the top solvers, and the per-solver design docs and CONFIG defaults.

References

- Sidky EY, Pan X. Image Reconstruction in Circular Cone-Beam Computed Tomography by Constrained, Total-Variation Minimization. *Physics in Medicine & Biology* 2008; 53(17): 4777–4807. doi: 10.1088/0031-9155/53/17/021
- Wu M, Maier A, Yang Q, Fahrig R. A Novel Filtered Backprojection-Based Algorithm for Sparse-View CT Image Reconstruction. In: ; 2015: 202–205.
- Manduca A, Yu L, Trzasko JD, et al. Projection Space Denoising with Bilateral Filtering and CT Noise Modeling for Dose Reduction in CT. *Medical Physics* 2009; 36(11): 4911–4919. doi: 10.1118/1.3232004
- Wagner F, Thies M, Gu M, et al. Ultralow-Parameter Denoising: Trainable Bilateral Filter Layers in Computed Tomography. *Medical Physics* 2022; 49(8): 5107–5120. doi: 10.1002/mp.15718
- Würfl T, Hoffmann M, Christlein V, et al. Deep Learning Computed Tomography: Learning Projection-Domain Weights From Image Domain in Limited Angle Problems. *IEEE Transactions on Medical Imaging* 2018; 37(6): 1454–1463. doi: 10.1109/TMI.2018.2833499
- Xu Y, Han S, Wang D, Wang G, Maltz JS, Yu H. Hybrid U-Net and Swin-Transformer Network for Limited-Angle Cardiac Computed Tomography. *Physics in Medicine & Biology* 2024; 69(10): 105012. doi: 10.1088/1361-6560/ad3db9
- Liang J, Cao J, Sun G, Zhang K, Van Gool L, Timofte R. SwinIR: Image Restoration Using Swin Transformer. In: ; 2021: 1833–1844. arXiv:2108.10257
- Liu Z, Lin Y, Cao Y, et al. Swin Transformer: Hierarchical Vision Transformer Using Shifted Windows. In: ; 2021: 9992–10002. arXiv:2103.14030
- Zhang Z, Yu L, Liang X, Zhao W, Xing L. TransCT: Dual-Path Transformer for Low-Dose Computed Tomography. In: Lecture Notes in Computer Science. Springer; 2021: 55–64. arXiv:2103.00634
- Hammernik K, Würfl T, Pock T, Maier A. A Deep Learning Architecture for Limited-Angle Computed Tomography Reconstruction. In: Informatik aktuell. Springer; 2017: 92–97
- Hammernik K, Klatzer T, Kobler E, et al. Learning a Variational Network for Reconstruction of Accelerated MRI Data. *Magnetic Resonance in Medicine* 2018; 79(6): 3055–3071. doi: 10.1002/mrm.26977
- Adler J, Öktem O. Learned Primal-Dual Reconstruction. *IEEE Transactions on Medical Imaging* 2018; 37(6): 1322–1332. doi: 10.1109/TMI.2018.2799231
- Wagner F, Thies M, Pfaff L, et al. On the Benefit of Dual-Domain Denoising in a Self-Supervised Low-Dose CT Setting. In: IEEE; 2023: 1–5
- Hendriksen AA, Pelt DM, Batenburg KJ. Noise2Inverse: Self-Supervised Deep Convolutional Denoising for Tomography. *IEEE Transactions on Computational Imaging* 2020; 6: 1320–1335. doi: 10.1109/TCI.2020.3019647
- Shi J, Pelt DM, Batenburg KJ. DM4CT: Benchmarking Diffusion Models for Computed Tomography Reconstruction. In: ; 2026. arXiv:2602.18589.
- Friedrich P, Wolleb J, Bieder F, Durrer A, Cattin PC. WDM: 3D Wavelet Diffusion Models for High-Resolution Medical Image Synthesis. In: LNCS. Springer; 2024. arXiv:2402.19043
- Song Y, Shen L, Xing L, Ermon S. Solving Inverse Problems in Medical Imaging with Score-Based Generative Models. In: ; 2022. arXiv:2111.08005.
- Chung H, Kim J, McCann MT, Klasky ML, Ye JC. Diffusion Posterior Sampling for General Noisy Inverse Problems. In: ; 2023. arXiv:2209.14687.
- Zha R, Zhang Y, Li H. NAF: Neural Attenuation Fields for Sparse-View CBCT Reconstruction. In: Lecture Notes in Computer Science. Springer; 2022: 442–452. arXiv:2209.14540
- Zha R, Lin TJ, Cai Y, Cao J, Zhang Y, Li H. R²-Gaussian: Rectifying Radiative Gaussian Splatting for Tomographic Reconstruction. In: ; 2024. arXiv:2405.20693.
- Liu H, Liu B, Hu Q, et al. A Review on 3D Gaussian Splatting for Sparse View Reconstruction. *Artificial Intelligence Review* 2025; 58: 215. doi: 10.1007/s10462-025-11171-4
- Terris M, Hurault S, Song M, Tachella J. Reconstruct Anything Model: A Lightweight Foundational Model for Computational Imaging. *arXiv preprint arXiv:2503.08915* 2025. arXiv:2503.08915.
- McCollough CH, Bartley AC, Carter RE, et al. Low-Dose CT for the Detection and Classification of Metastatic Liver Lesions: Results of the 2016 Low Dose CT Grand Challenge. *Medical Physics* 2017; 44(10): e339–e352. doi: 10.1002/mp.12345
- Noo F, Defrise M, Clackdoyle R. Single-Slice Rebinning Method for Helical Cone-Beam CT. *Physics in Medicine & Biology* 1999; 44(2): 561–570. doi: 10.1088/0031-9155/44/2/019
- Syben C, Michen M, Stimpel B, Seitz S, Ploner S, Maier AK. Technical Note: PYRO-NN: Python Reconstruction Operators in Neural Networks. *Medical Physics* 2019; 46(11): 5110–5115. doi: 10.1002/mp.13753
- Sidky EY, Pan X. Report on the AAPM Deep-Learning Sparse-View CT (DL-Sparse-View CT) Grand Challenge. *Medical Physics* 2022; 49(8): 4935–4943. doi: 10.1002/mp.15489
- Sidky EY, Pan X. Report on the AAPM Deep-Learning Spectral CT Grand Challenge. *Medical Physics* 2024. arXiv:2212.06718doi: 10.1002/mp.16612
- Abadi E, Segars WP, Felice N, Hoffman EA, Samei E. The AAPM Truth-Based CT (TrueCT) Reconstruction Grand Challenge. *Medical Physics* 2025; 52(4): 1978–1990. doi: 10.1002/mp.17619
- Haneda E, Peters N, Zhang J, et al. AAPM CT Metal Artifact Reduction Grand Challenge. *Medical Physics* 2025; 52(10): e70050. doi: 10.1002/mp.70050
- Maier A, Yang SH, Maleki F, Muthukrishnan N, Forghani R. Offer Proprietary Algorithms Still Protection of Intellectual Property in the Age of Machine Learning? A Case Study Using Dual Energy CT Data. In: Informatik aktuell. Springer Vieweg, Wiesbaden. 2022 (pp. 345–350). arXiv:2112.03678
- Maier AK, Syben C, Stimpel B, et al. Learning with Known Operators Reduces Maximum Error Bounds. *Nature Machine Intelligence* 2019; 1(8): 373–380. doi: 10.1038/s42256-019-0077-5
- Maier A, Köstler H, Heisig M, Krauss P, Yang SH. Known Operator Learning and Hybrid Machine Learning in Medical Imaging—A Review of the Past, the Present, and the Future. *Progress in Biomedical Engineering* 2022; 4(2): 022002. doi: 10.1088/2516-1091/ac5b13
- Maier A, Syben C, Lasser T, Riess C. A Gentle Introduction to Deep Learning in Medical Image Processing. *Zeitschrift für Medizinische Physik* 2019; 29(2): 86–101. doi: 10.1016/j.zemedi.2018.12.003
- Ongie G, Jalal A, Metzler CA, Baraniuk RG, Dimakis AG, Willett R. Deep Learning Techniques for Inverse Problems in Imaging. *IEEE Journal on Selected Areas in Information Theory* 2020; 1(1): 39–56. doi: 10.1109/JSAIT.2020.2991563

35. Lu C, Lu C, Lange RT, Foerster J, Clune J, Ha D. The AI Scientist: Towards Fully Automated Open-Ended Scientific Discovery. *arXiv preprint arXiv:2408.06292* 2024. arXiv:2408.06292.
36. Yang J, Jimenez CE, Wettig A, et al. SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering. In: ; 2024. arXiv:2405.15793.
37. Jimenez CE, Yang J, Wettig A, et al. SWE-bench: Can Language Models Resolve Real-World GitHub Issues?. In: ; 2024. arXiv:2310.06770.
38. Zaiss M, Aly A, Endres J, Dornstetter T, Weinmüller S, Maier A. MR LLMs with MR Agentic Sequence Development: Leveraging Skills for Automatic Physics-Informed Sequence Development. *arXiv preprint arXiv:2604.13282* 2026. arXiv:2604.13282; Agent4MR.